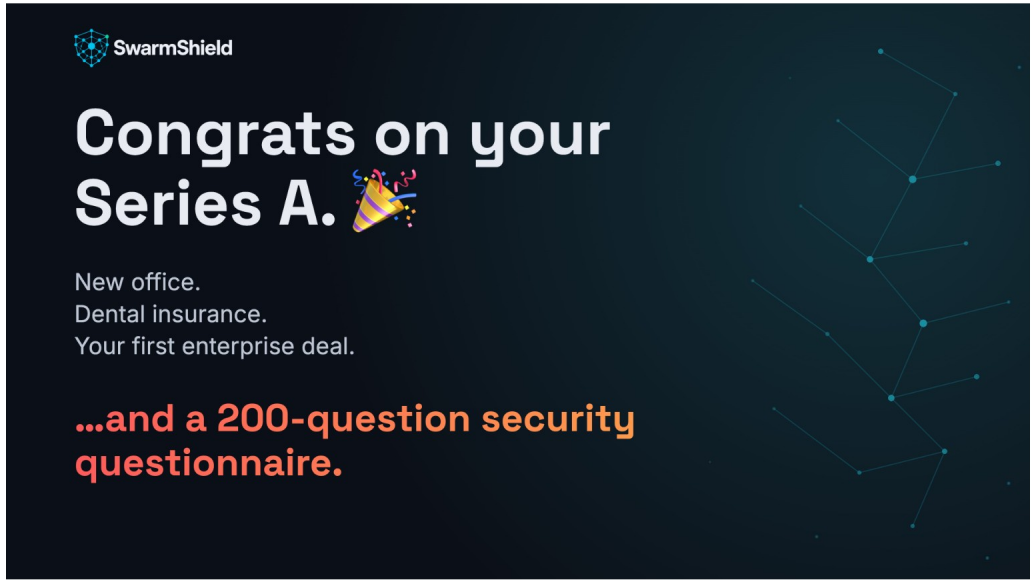


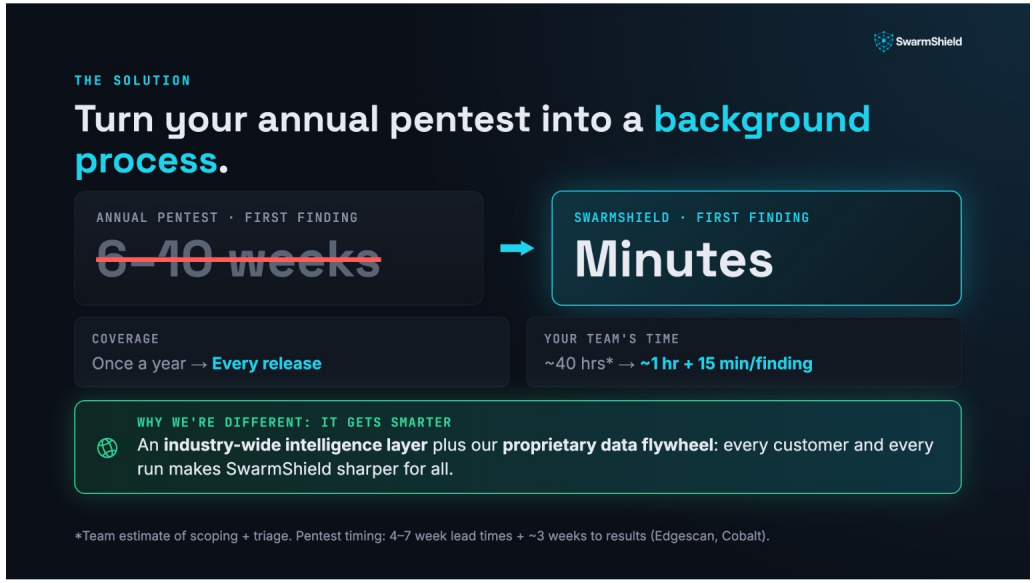
Congrats — you just raised your Series A. Bigger office, real benefits, and your first enterprise customer. And then their security team sends you a 200-question questionnaire. That's the moment we built SwarmShield for.



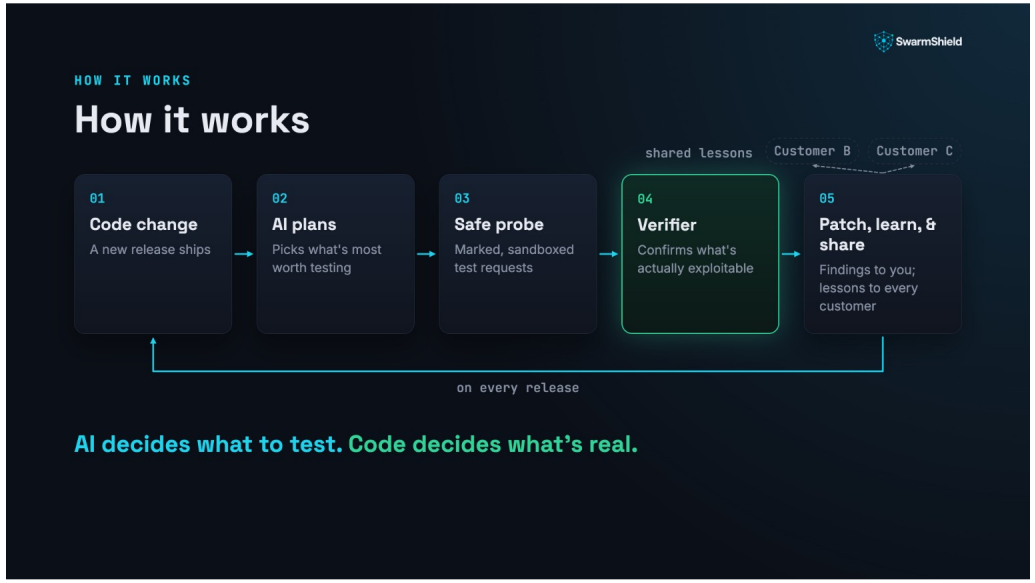
Your first enterprise deal is usually the first time you have to take security seriously — 85 to 95 percent of enterprise buyers require SOC 2. So you hire one security person. Realistically, they can run one pentest a year: it takes six to ten weeks from booking to report and costs around eighteen thousand dollars. Meanwhile your engineers ship every day, increasingly with AI — and Veracode found 45% of AI-generated code has security flaws.



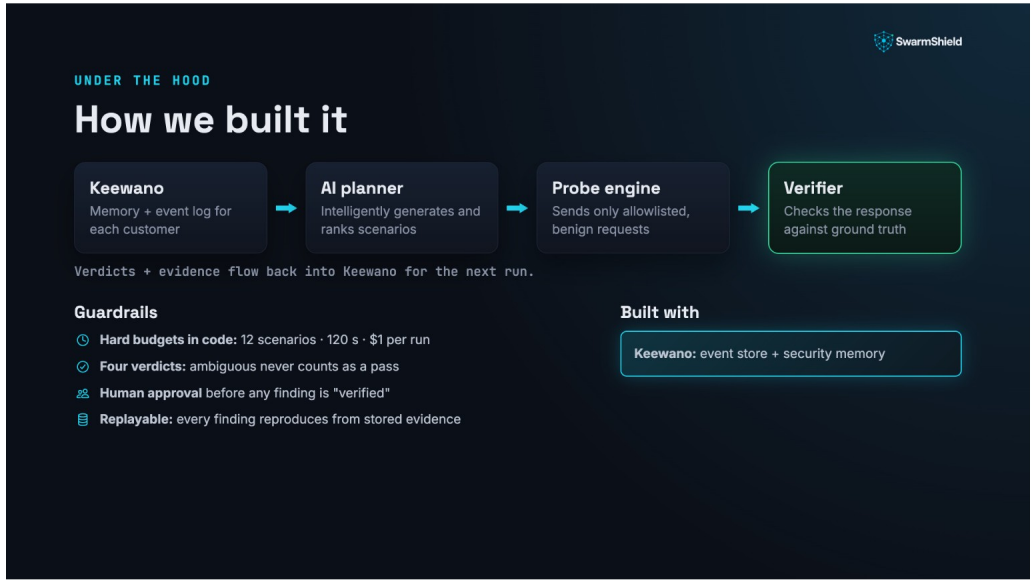
Attackers move in hours — Verizon found the median time from a vulnerability's disclosure to mass exploitation is zero days. Defenders test once a year — about half of companies pentest once a year or less. That gap is the problem we're solving.



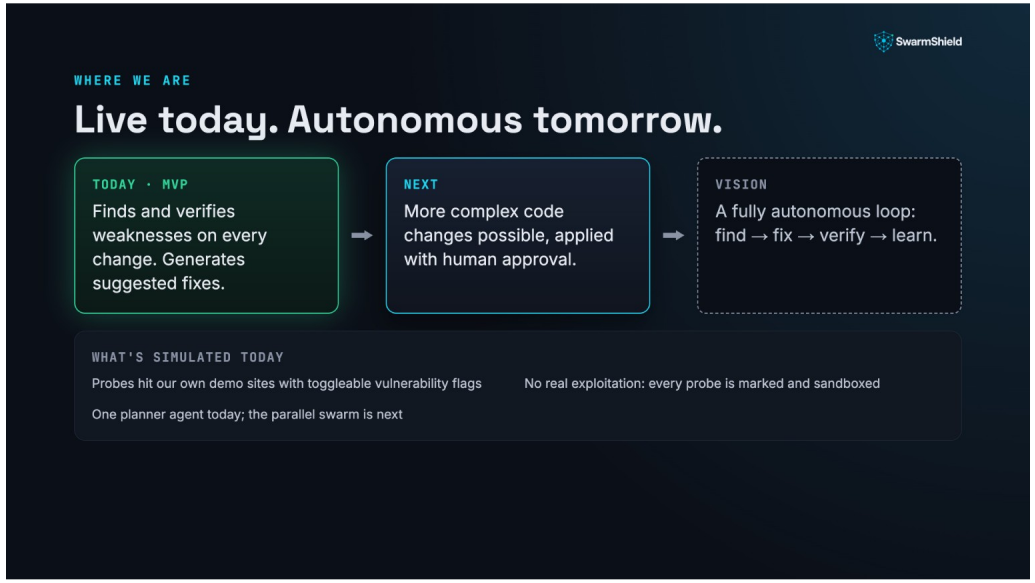
SwarmShield turns that once-a-year pentest into a background process. Every time you ship, it decides what changed, safely tests it the way an attacker would, and verifies whether the weakness is real. Your first finding arrives minutes after a deploy instead of six to ten weeks later — and your security person spends about fifteen minutes reviewing each finding instead of a week scoping and triaging a report. That includes setup and review time — we're not hiding the human part. And here's why we're different: we're building an industry-wide intelligence layer. Every verified finding, scrubbed of anything that identifies the customer, becomes a lesson that tests every other customer. Combined with our proprietary data flywheel — every run's events, verdicts, and evidence — SwarmShield gets smarter over time. A pentest firm starts from zero on every engagement. We never do.



Here's the loop. A code change ships. An AI planner reads what changed and what we already know about this customer, and picks the tests most worth running. A probe engine sends safe, clearly marked requests. Then — and this is the important part — a deterministic verifier, not the AI, decides whether the weakness is real. Findings go to you, and generalized, privacy-scrubbed lessons go to every other customer. AI decides what to test. Code decides what's real.



Under the hood: Keewano stores every event and each customer's security memory. An AI planner ranks scenarios from a fixed catalog — it returns scenario IDs, not code. A probe engine executes only allowlisted, benign requests, and a deterministic verifier compares the response against ground truth. Budgets are enforced in code, not in the prompt, and nothing becomes "verified" without a human. Every confirmed finding can be replayed from its stored evidence. [Add one sentence per sponsor actually used.]



Here's what's real today: SwarmShield finds and verifies weaknesses on every change and hands your team an evidence-backed report. To be transparent — today it runs against our own demo sites with simulated vulnerabilities, and it's a single planner agent rather than a full parallel swarm. Next, it opens fix PRs that a human approves. The vision is a fully autonomous, self-healing loop: find, fix, verify, and learn.



Attackers move in hours. With SwarmShield, so do defenders. A security team of one gets a swarm. Thank you.